



**BATCH** : B107 AWS-DevOps  
**LESSON** : **Kubernetes**  
**DATE** : 20.05.2023  
**SUBJECT** : **Kubernetes Objects**

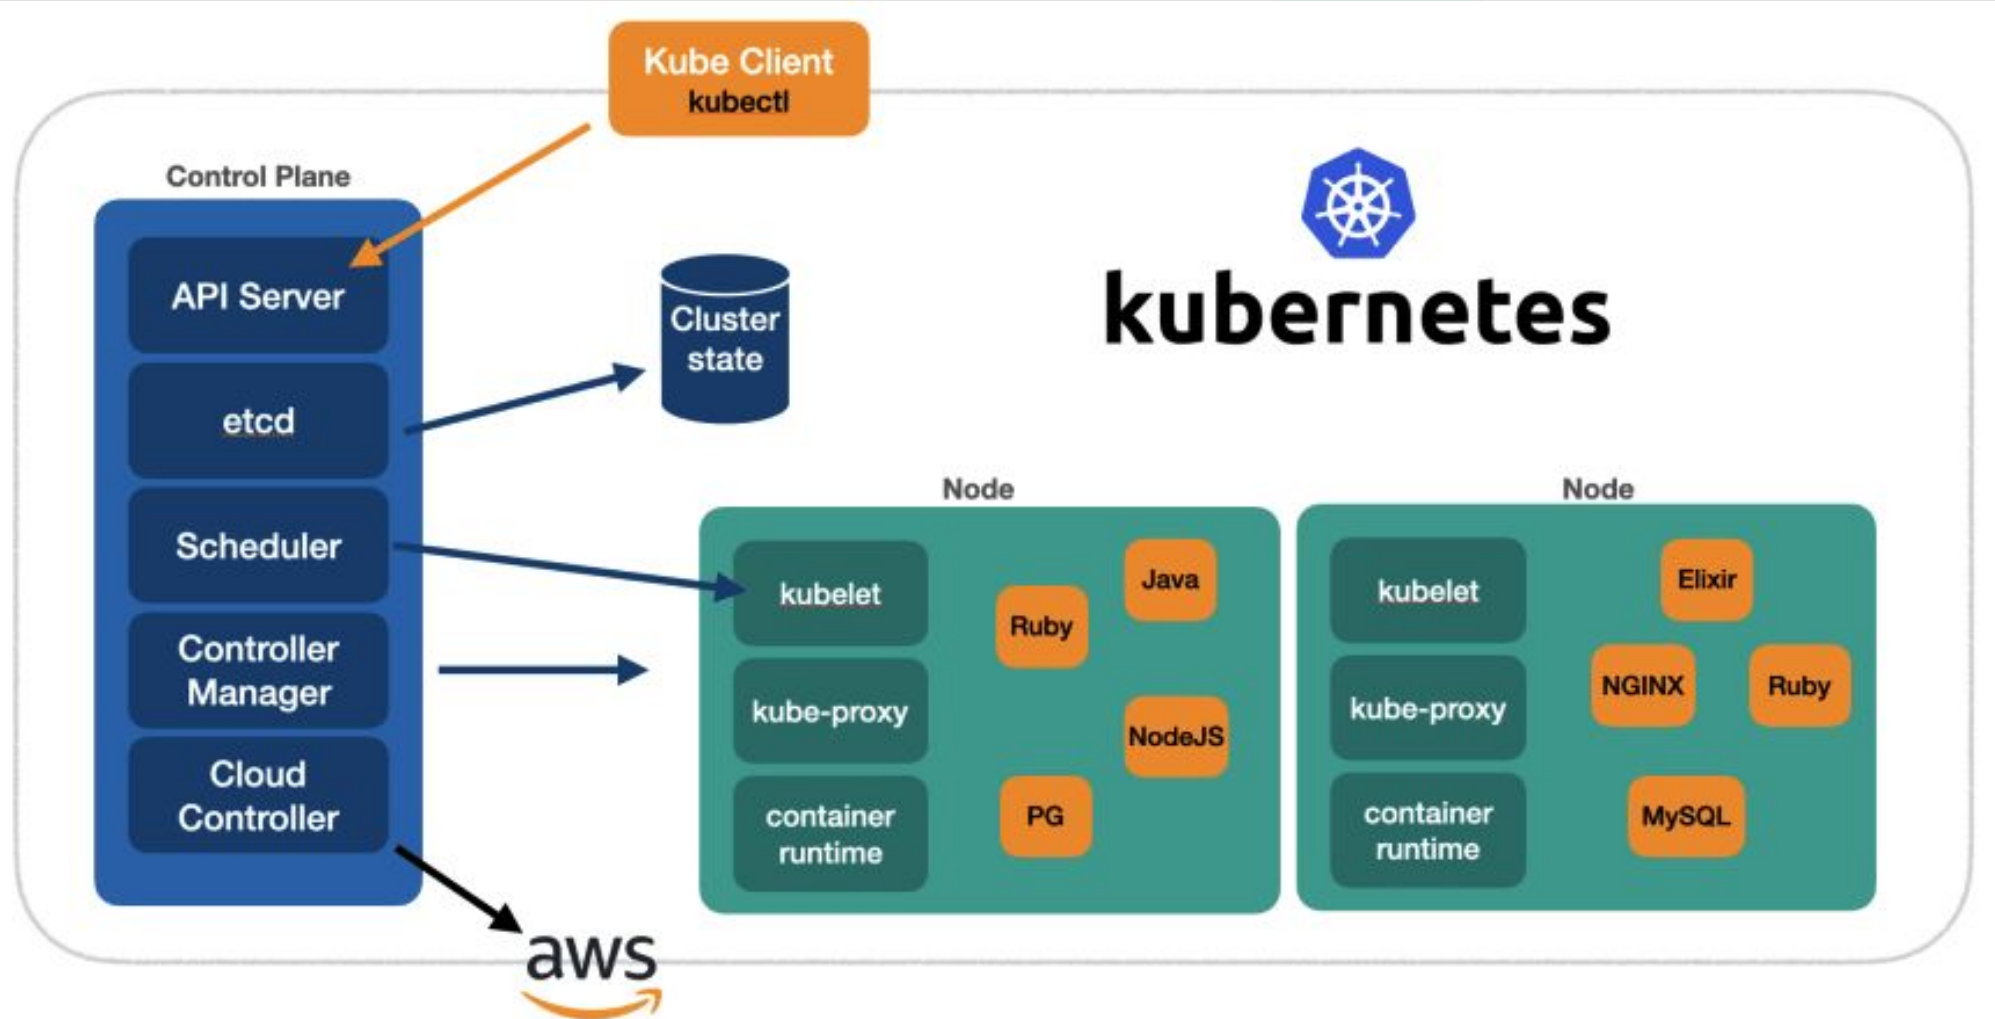
ZOOM GİRİŞLERİNİZİ LÜTFEN **LMS** SİSTEMİ ÜZERİNDEN YAPINIZ





Kubernetes







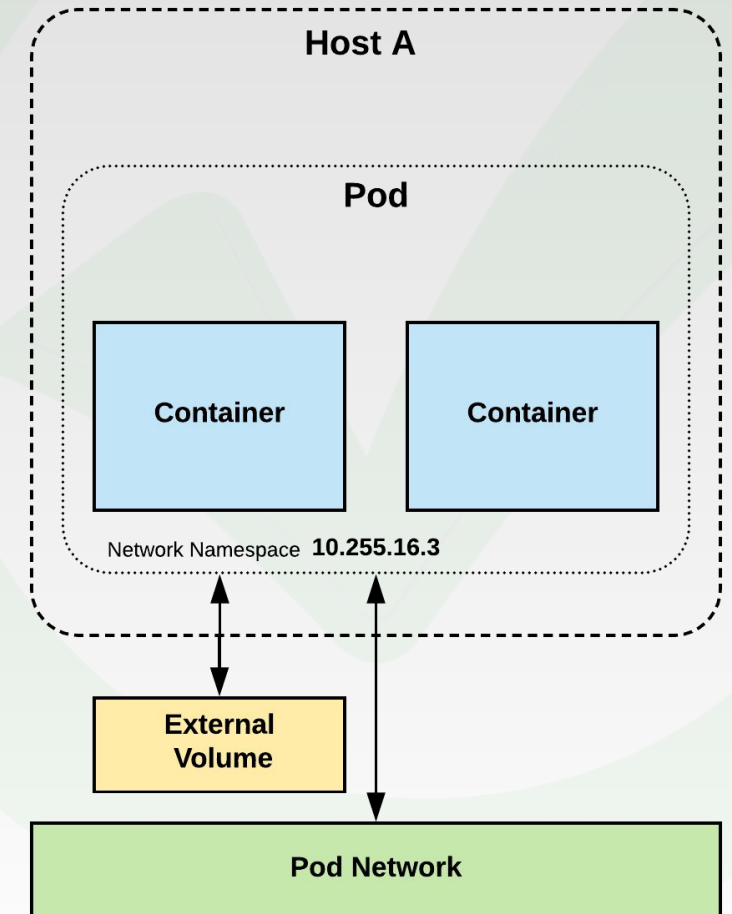
# Important Terms

- In robotics and automation, **a control loop** is a non-terminating loop that regulates the state of a system. Thermostat in a room can be an example. When you set the temperature, that's telling the thermostat about your **desired state**. The actual room temperature is the **current state**.
- In Kubernetes, **controllers are control loops that watch the state of your cluster**. Each controller tries to move the current cluster state closer to the desired state.
- **A workload is an application** running on Kubernetes. Whether your workload is **a single component or several** that work together, on Kubernetes you run it inside a set of pods.
- You can use **workload resources** that manage a set of pods on your behalf. These **resources configure controllers** that make sure the right number of the right kind of pod are running, to match the state you specified.
- This **desired state** is mentioned in the **spec field**.



# Pods

- Atomic unit or **smallest “unit of work”** of Kubernetes.
- Pods contain **one or MORE containers** that share volumes, a network, and are a part of a single context.
- Pods are used in two main ways:
  - Pods that run a single container
  - Pods that run multiple containers that need to work together
- Pods are virtual machines.
- **Nodes are physical machines.**







# Pod Template

- Workload Controllers manage instances of Pods based off a provided template.
- Pod Templates are Pod specs with limited metadata.
- Controllers use Pod Templates to make actual pods.

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx
```

```
template:
  metadata:
    labels:
      app: nginx
  spec:
    containers:
    - name: nginx
      image: nginx
```



# ReplicaSet

- Primary method of **managing pod replicas and their lifecycle.**
- Includes their **scheduling, scaling, and deletion.**
- Their job is simple: **Always ensure the desired number of pods are running.**





# ReplicaSet

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-example
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  template:
    <pod template>
```

- **replicas**: The desired number of instances of the Pod.
- **selector**: The label selector for the ReplicaSet will manage ALL Pod instances that it targets; whether it's desired or not.

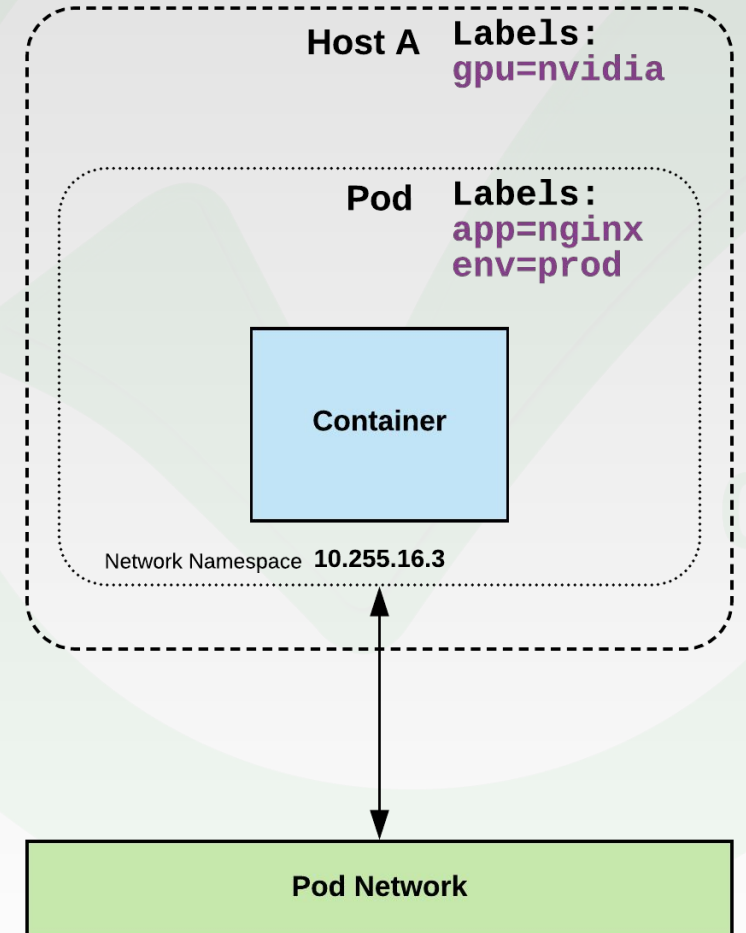




# Labels

- **key-value pairs** that are used to **identify, describe and group** together related sets of objects or resources.
- pods and nodes can be labeled
- NOT characteristic of uniqueness.
- Have a strict syntax with a slightly limited character set.

\* <https://kubernetes.io/docs/concepts/overview/working-with-objects/labels/#syntax-and-character-set>






# Label Example

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-label-example
  labels:
    app: nginx
    env: prod
spec:
  containers:
  - name: nginx
    image: nginx:stable-alpine
    ports:
    - containerPort: 80
```

```
kubectl label pod nginx app=nginx
kubectl label node host1 gpu=nvidia
```





# Selectors

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-label-example
  labels:
    app: nginx
    env: prod
spec:
  containers:
    - name: nginx
      image:
nginx:stable-alpine
      ports:
        - containerPort: 80
  nodeSelector:
    gpu: nvidia
```

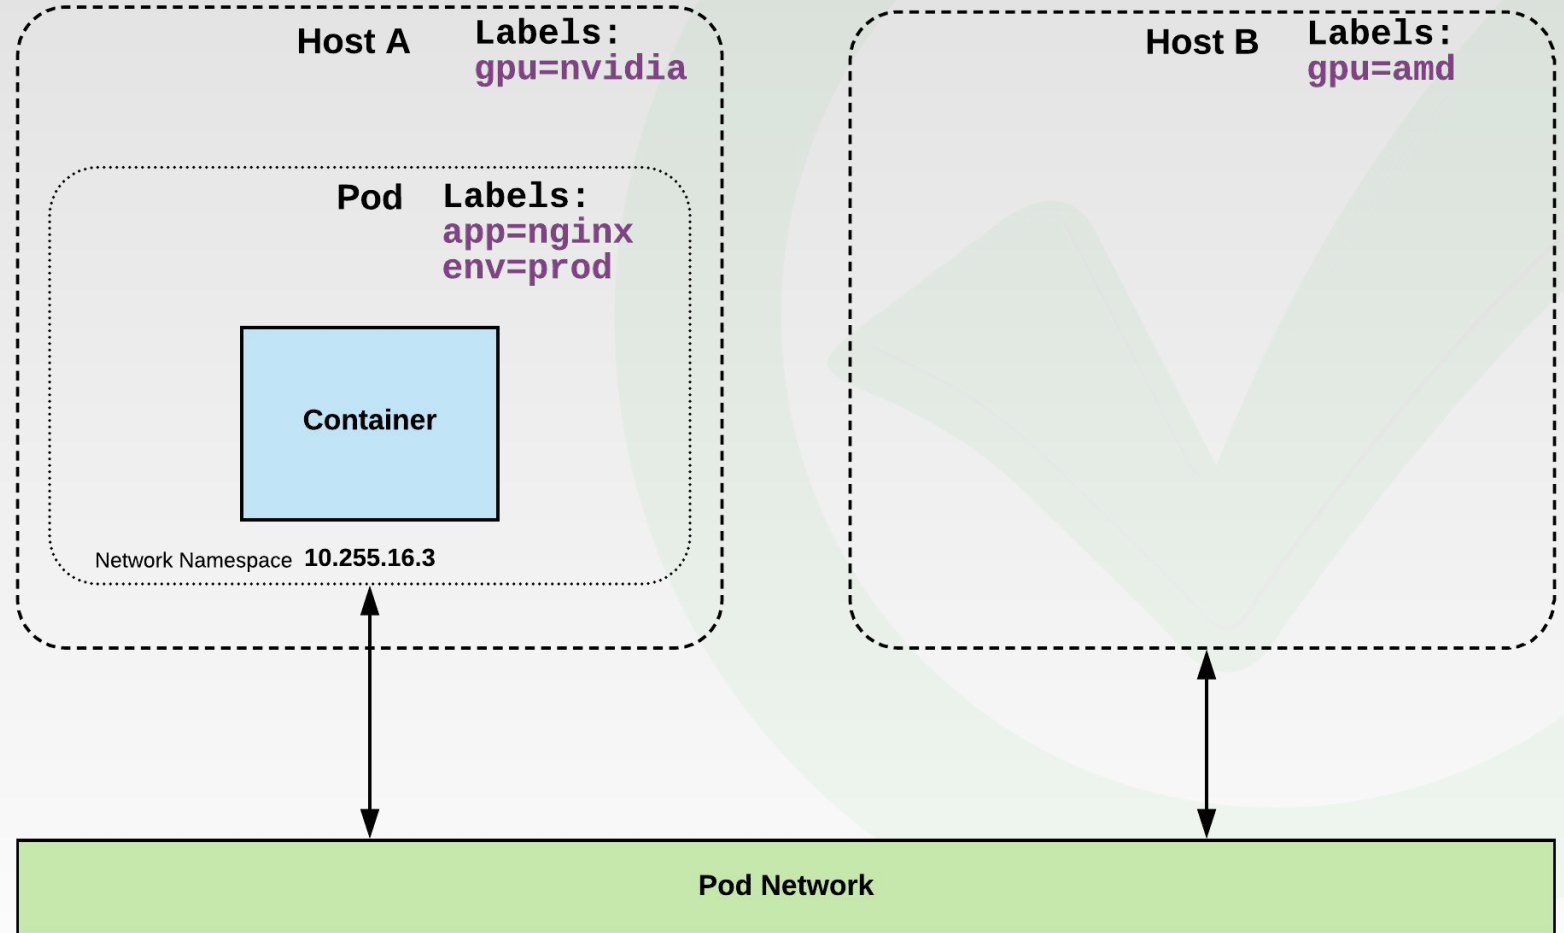
Selectors use labels to filter or select objects, and are used throughout Kubernetes.

```
selector:
  matchLabels:
    app: nginx
```



# Selector Example

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-label-example
  labels:
    app: nginx
    env: prod
spec:
  containers:
  - name: nginx
    image: nginx:stable-alpine
    ports:
    - containerPort: 80
  nodeSelector:
    gpu: nvidia
```





# Deployment

- Declarative method of **managing Pods via ReplicaSets.**
- Provide **rollback functionality and update control.**
- Updates are managed through the pod-template-hash label.
- Each iteration creates a unique label that is assigned to both the ReplicaSet and subsequent Pods.





# Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: deploy-example
spec:
  replicas: 3
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  template:
    <pod template>
```

- **revisionHistoryLimit:** The number of previous iterations of the Deployment to retain.
- **strategy:** Describes the method of updating the Pods based on the **type**. Valid options are **Recreate** or **RollingUpdate**.
  - **Recreate:** All existing Pods are killed before the new ones are created.
  - **RollingUpdate:** Cycles through updating the Pods according to the parameters: **maxSurge** and **maxUnavailable**.
    - **maxUnavailable:** how many pods can be unavailable at most during update
    - **maxSurge:** how many pods can exceed the desired number of Pods during update





# Namespaces

```
apiVersion: v1
kind: Namespace
metadata:
  name: prod
  labels:
    app: MyBigWebApp
```

- Namespaces are a logical cluster or environment, and are the primary method of partitioning a cluster or scoping access.
- Namespaces are a way to divide cluster resources between multiple users.

```
$ kubectl get ns --show-labels
```

| NAME        | STATUS | AGE | LABELS          |
|-------------|--------|-----|-----------------|
| default     | Active | 11h | <none>          |
| kube-public | Active | 11h | <none>          |
| kube-system | Active | 11h | <none>          |
| prod        | Active | 6s  | app=MyBigWebApp |



# Rolling Update Deployment

Updating pod template generates a new **ReplicaSet** revision.

R1 pod-template-hash:

676677fff

R2 pod-template-hash:

54f7ff7d6d

Deployment  
Revision 1

ReplicaSet R1

ReplicaSet R2

Pod

Pod

Pod

```
$ kubectl get replicaset
```

| NAME            | DESIRED | CURRENT | READY | AGE |
|-----------------|---------|---------|-------|-----|
| mydep-676677fff | 3       | 3       | 3     | 5h  |

```
$ kubectl get pods
```

| NAME                  | READY | STATUS  | RESTARTS | AGE |
|-----------------------|-------|---------|----------|-----|
| mydep-676677fff-9r2zn | 1/1   | Running | 0        | 5h  |
| mydep-676677fff-hsfz9 | 1/1   | Running | 0        | 5h  |
| mydep-676677fff-sjxhf | 1/1   | Running | 0        | 5h  |



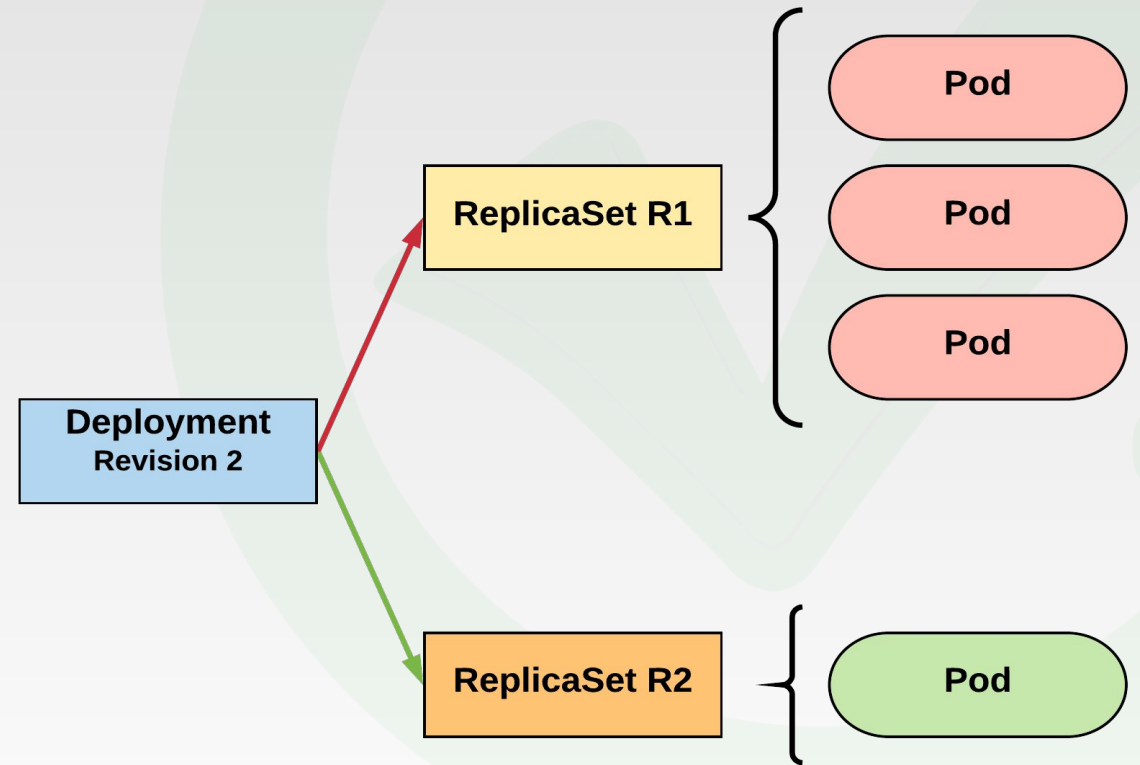
# Rolling Update Deployment

New **ReplicaSet** is initially scaled up based on **maxSurge**.

R1 pod-template-hash:  
**676677fff**  
R2 pod-template-hash:  
**54f7ff7d6d**

```
$ kubectl get replicaset
NAME                                DESIRED  CURRENT  READY  AGE
mydep-54f7ff7d6d                    1         1        1      5s
mydep-6766777fff                    2         3        3      5h
```

```
$ kubectl get pods
NAME                                READY    STATUS    RESTARTS  AGE
mydep-54f7ff7d6d-9gv1l             1/1      Running   0          2s
mydep-6766777fff-9r2zn             1/1      Running   0          5h
mydep-6766777fff-hsfz9             1/1      Running   0          5h
mydep-6766777fff-sjxhf             1/1      Running   0          5h
```





# Rolling Update Deployment

Phase out of old Pods managed by `maxSurge` and `maxUnavailable`.

R1 pod-template-hash:  
**676677fff**  
R2 pod-template-hash:  
**54f7ff7d6d**

Deployment  
Revision 2

ReplicaSet R1

Pod

Pod

~~Pod~~

ReplicaSet R2

Pod

Pod

```
$ kubectl get replicaset
```

| NAME                    | DESIRED  | CURRENT  | READY    | AGE       |
|-------------------------|----------|----------|----------|-----------|
| <b>mydep-54f7ff7d6d</b> | <b>2</b> | <b>2</b> | <b>2</b> | <b>8s</b> |
| mydep-676677fff         | 2        | 2        | 2        | 5h        |

```
$ kubectl get pods
```

| NAME                          | READY      | STATUS         | RESTARTS | AGE       |
|-------------------------------|------------|----------------|----------|-----------|
| <b>mydep-54f7ff7d6d-9gv1l</b> | <b>1/1</b> | <b>Running</b> | <b>0</b> | <b>5s</b> |
| <b>mydep-54f7ff7d6d-cqv1q</b> | <b>1/1</b> | <b>Running</b> | <b>0</b> | <b>2s</b> |
| mydep-676677fff-9r2zn         | 1/1        | Running        | 0        | 5h        |
| mydep-676677fff-hsfz9         | 1/1        | Running        | 0        | 5h        |



# Rolling Update Deployment

Phase out of old Pods managed by `maxSurge` and `maxUnavailable`.

R1 pod-template-hash:  
**676677fff**  
R2 pod-template-hash:  
**54f7ff7d6d**

Deployment  
Revision 2

```
$ kubectl get replicaset
```

| NAME             | DESIRED | CURRENT | READY | AGE |
|------------------|---------|---------|-------|-----|
| mydep-54f7ff7d6d | 3       | 3       | 3     | 10s |
| mydep-6766777fff | 0       | 1       | 1     | 5h  |

```
$ kubectl get pods
```

| NAME                   | READY | STATUS  | RESTARTS | AGE |
|------------------------|-------|---------|----------|-----|
| mydep-54f7ff7d6d-9gv1l | 1/1   | Running | 0        | 7s  |
| mydep-54f7ff7d6d-cqvlq | 1/1   | Running | 0        | 5s  |
| mydep-54f7ff7d6d-gccr6 | 1/1   | Running | 0        | 2s  |
| mydep-6766777fff-9r2zn | 1/1   | Running | 0        | 5h  |

ReplicaSet R1

ReplicaSet R2

Pod

~~Pod~~

~~Pod~~

Pod

Pod

Pod



# Rolling Update Deployment

Phase out of old Pods managed by  
**maxSurge** and **maxUnavailable**.

R1 pod-template-hash:

676677fff

R2 pod-template-hash:

54f7ff7d6d

```
$ kubectl get replicaset
```

| NAME             | DESIRED | CURRENT | READY | AGE |
|------------------|---------|---------|-------|-----|
| mydep-54f7ff7d6d | 3       | 3       | 3     | 13s |
| mydep-6766777fff | 0       | 0       | 0     | 5h  |

```
$ kubectl get pods
```

| NAME                   | READY | STATUS  | RESTARTS | AGE |
|------------------------|-------|---------|----------|-----|
| mydep-54f7ff7d6d-9gv1l | 1/1   | Running | 0        | 10s |
| mydep-54f7ff7d6d-cqv1q | 1/1   | Running | 0        | 8s  |
| mydep-54f7ff7d6d-gccr6 | 1/1   | Running | 0        | 5s  |

Deployment  
Revision 2

ReplicaSet R1

ReplicaSet R2

~~Pod~~

~~Pod~~

~~Pod~~

Pod

Pod

Pod





# RollingUpdate Deployment

Updated to new deployment revision completed.

R1 pod-template-hash:

**676677fff**

R2 pod-template-hash:

**54f7ff7d6d**

Deployment  
Revision 2

ReplicaSet R1

ReplicaSet R2

Pod

Pod


Pod

```
$ kubectl get replicaset
```

| NAME             | DESIRED | CURRENT | READY | AGE |
|------------------|---------|---------|-------|-----|
| mydep-54f7ff7d6d | 3       | 3       | 3     | 15s |
| mydep-6766777fff | 0       | 0       | 0     | 5h  |

```
$ kubectl get pods
```

| NAME                   | READY | STATUS  | RESTARTS | AGE |
|------------------------|-------|---------|----------|-----|
| mydep-54f7ff7d6d-9gv1l | 1/1   | Running | 0        | 12s |
| mydep-54f7ff7d6d-cqvlq | 1/1   | Running | 0        | 10s |
| mydep-54f7ff7d6d-gccr6 | 1/1   | Running | 0        | 7s  |

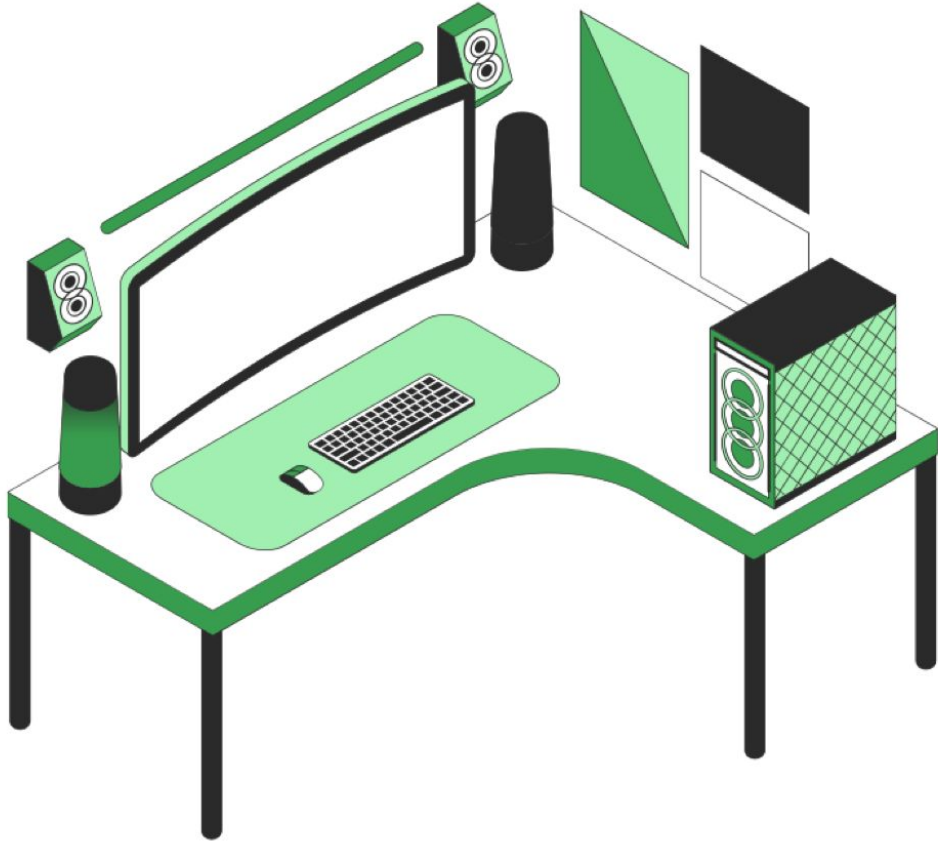




# Summary

## In a K8s cluster:

- I need to start a containerized app. Solution: use .. Pod
- I need to launch a containerized app of 10 replicas. Solution: use .. ReplicaSet
- I need to configure updates on a containerized app. Solution: use .. Deployment
- I need to make sure a specific group of Pods are run on a specific node. Labels, Selectors
- I need to separate resources of Java Dev from the rest of the team. Namespaces



Do you  
have any  
questions?

Send it to us! We hope you learned  
something new.